

Exploratory Data Analysis of Spotify Songs

Nama: Banu Agil Triyanto

NIM: 18225109

Dataset: Spotify Songs (TidyTuesday, 2020-01-21), berisi audio features untuk ~33 ribu lagu dari 6 genre.

Tujuan: melakukan EDA yang lengkap dan reproducible menggunakan **hanya NumPy dan Pandas** (tanpa scipy / sklearn di bagian yang dilarang), menjawab pertanyaan analitis yang diminta, dan mendokumentasikan setiap keputusan.

Alur kerja AI tools (bagaimana notebook ini dibuat): Saya memakai AI coding assistant sebagai *pair programmer*, bukan autopilot. Konkretnya:

- **Gemini Code Assist / GitHub Copilot** untuk membantu menulis docstring dan mengecek ulang NumPy vectorized saya (Stage 4, Bonus 3).
- **Claude Code** untuk membantu menyusun struktur notebook dan menguji edge case (penanganan duplikat, pembagian nol pada min-max scaling).

Setiap saran AI saya review, jalankan, lalu saya acc, edit, atau delete. Prosesnya saya dokumentasikan di Stage 4.

Reproducibility: notebook ini jalan dari atas ke bawah tanpa setup manual. Yang dibutuhkan hanya `pandas` dan `numpy`.

Stage 1: Setup and Initial Inspection

Kita load data langsung dari source URL supaya notebook sepenuhnya self-contained.

```
In [1]: import numpy as np
import pandas as pd

pd.set_option("display.max_columns", None)
pd.set_option("display.width", 140)
np.set_printoptions(precision=4, suppress=True)

print("pandas", pd.__version__, "| numpy", np.__version__)
```

pandas 3.0.3 | numpy 2.5.0

1.1 Load dataset

```
In [2]: URL = ("https://raw.githubusercontent.com/rfordatascience/tidytuesday/"
              "master/data/2020/2020-01-21/spotify_songs.csv")

df_raw = pd.read_csv(URL)
print("Loaded:", df_raw.shape[0], "rows x", df_raw.shape[1], "columns")
df_raw.head(3)
```

Loaded: 32833 rows x 23 columns

```
Out[2]:
```

	track_id	track_name	track_artist	track_popularity	tr
0	6f807x0ima9a1j3VPbc7VN	I Don't Care (with Justin Bieber) - Loud Luxur...	Ed Sheeran	66	2oCs0DGTsRk
1	0r7CVbZTWZgbTCYdfa2P31	Memories - Dillon Francis Remix	Maroon 5	67	63rPSO264uF
2	1z1Hg7Vb0AhHDEmnDE79I	All the Time - Don Diablo Remix	Zara Larsson	70	1HoSmj2eLc

1.2 Inspeksi awal: shape, dtypes, missing values

```
In [3]: print("Shape:", df_raw.shape, "\n")
print("Dtypes:")
print(df_raw.dtypes)
```

Shape: (32833, 23)

```
Dtypes:
track_id          str
track_name        str
track_artist      str
track_popularity  int64
track_album_id   str
track_album_name  str
track_album_release_date  str
playlist_name     str
playlist_id       str
playlist_genre    str
playlist_subgenre str
danceability      float64
energy            float64
key              int64
loudness          float64
mode             int64
speechiness       float64
acousticness      float64
instrumentalness  float64
liveness          float64
valence           float64
tempo            float64
duration_ms       int64
dtype: object
```

```
In [4]: missing = df_raw.isna().sum()
missing_pct = (missing / len(df_raw) * 100).round(3)
missing_table = (pd.DataFrame({"n_missing": missing, "pct_missing": missing_
                             .query("n_missing > 0")
                             .sort_values("n_missing", ascending=False)})
print("Columns with missing values:")
missing_table
```

Columns with missing values:

```
Out[4]:
```

	n_missing	pct_missing
track_name	5	0.015
track_artist	5	0.015
track_album_name	5	0.015

Menangani aturan missing >20%. Hanya `track_name`, `track_artist`, dan `track_album_name` yang punya missing value, dan masing-masing cuma **5 baris** (~0.015%), jauh di bawah threshold 20%. Jadi **tidak ada kolom yang di-drop**.

Kelima baris itu tidak punya nama artist, yang kita butuhkan untuk analisis level artist di Stage 2, jadi kita drop *baris-baris itu* (bukan kolomnya). Kehilangan ini sangat kecil (5 dari 32.833).

```
In [5]: cols_over_20 = missing_pct[missing_pct > 20].index.tolist()
print("Columns exceeding 20% missing:", cols_over_20 if cols_over_20 else "None")

# Drop the 5 rows with no artist/track metadata (needed downstream).
df = df_raw.dropna(subset=["track_artist"]).copy()
print("Rows after dropping metadata-less rows:", len(df))
```

Columns exceeding 20% missing: None
 Rows after dropping metadata-less rows: 32828

1.3 Penanganan duplikat

```
In [6]: print("Fully-identical rows:", df.duplicated().sum())
print("Rows sharing a track_id:", df.duplicated(subset='track_id').sum())

# Inspect one example of a repeated track_id.
dup_id = df[df.duplicated(subset='track_id', keep=False)][['track_id']].iloc[0]
df[df.track_id == dup_id][['track_id', 'track_name', 'playlist_genre', 'playlist_name']]
```

Fully-identical rows: 0
 Rows sharing a track_id: 4476

```
Out [6]:
```

	track_id	track_name	playlist_genre	playlist_name
0	6f807x0ima9a1j3VPbc7VN	I Don't Care (with Justin Bieber) - Loud Luxur...	pop	Pop Remix
29684	6f807x0ima9a1j3VPbc7VN	I Don't Care (with Justin Bieber) - Loud Luxur...	edm	Pop EDM Remixes

Strategi deduplikasi: dedup berdasarkan `track_id`.

Tidak ada baris yang benar-benar identik, tapi ada ~4.476 baris yang berbagi `track_id` yang sama. Penyebabnya struktural: *unit dari tabel mentah adalah pasangan ($track \times playlist$)*, jadi satu lagu muncul sekali untuk tiap playlist yang memuatnya, kadang di genre yang berbeda.

Untuk analisis level lagu (audio features, popularity, output artist), duplikasi playlist ini menghitung lagu berkali-kali dan membuat setiap mean jadi bias. `track_id` adalah identifier lagu yang unik secara global di Spotify, jadi itulah key yang benar. Kita simpan kemunculan pertama.

Trade-off: sebuah lagu yang ada di playlist `pop` dan `rock` akan diberi label genre yang muncul lebih dulu. Label genre di sini berasal dari playlist, bukan ground truth, jadi kita menerima sedikit noise label sebagai ganti satu baris per lagu.

```
In [7]: df = df.drop_duplicates(subset='track_id', keep='first').reset_index(drop=True)
print("Final analysis frame:", df.shape)
```

Final analysis frame: (28352, 23)

1.4 Tabel ringkasan numerik (IQR dihitung manual dengan NumPy)

```
In [8]: def numeric_summary(frame: pd.DataFrame) -> pd.DataFrame:
        """Calculate summary statistics (mean, median, std, and IQR) for numeric
        columns in a DataFrame.

        Parameters
        -----
        frame : pd.DataFrame
            The input DataFrame containing the data to summarize.

        Returns
        -----
        pd.DataFrame
            A summary DataFrame with columns for 'mean', 'median', 'std',
            and 'IQR', indexed by the original numeric column names. Values
            are rounded to three decimal places.

        Notes
        -----
        Non-numeric columns are automatically ignored.
        The interquartile range (IQR) is calculated manually using `np.percentile`
        with linear interpolation.
        """
        num = frame.select_dtypes(include=np.number)
        rows = {}
        for col in num.columns:
            vals = num[col].to_numpy()
            q1, q3 = np.percentile(vals, [25, 75])  # manual IQR
            rows[col] = {
                "mean": np.mean(vals),
                "median": np.median(vals),
                "std": np.std(vals, ddof=1),
                "IQR": q3 - q1,
            }
        return pd.DataFrame(rows).T.round(3)

summary = numeric_summary(df)
summary
```

Out [8]:

	mean	median	std	IQR
track_popularity	39.335	42.000	23.699	37.000
danceability	0.653	0.670	0.146	0.199
energy	0.698	0.722	0.184	0.264
key	5.367	6.000	3.614	7.000
loudness	-6.818	-6.261	3.036	3.601
mode	0.566	1.000	0.496	1.000
speechiness	0.108	0.063	0.103	0.092
acousticness	0.177	0.080	0.223	0.246
instrumentalness	0.091	0.000	0.233	0.007
liveness	0.191	0.127	0.156	0.156
valence	0.510	0.512	0.234	0.366
tempo	120.958	121.994	26.955	34.027
duration_ms	226574.631	216933.000	61081.364	67234.000

Stage 2: Genre and Popularity Analysis

2.5 Statistik distribusi per genre

```
In [9]: targets = ["track_popularity", "danceability", "energy", "valence"]
genre_stats = df.groupby("playlist_genre")[targets].agg(["mean", "std", "median"])
genre_stats
```

Out [9]:

	track_popularity			danceability			energy		
	mean	std	median	mean	std	median	mean	std	median
playlist_genre									
edm	30.678	20.347	33.0	0.658	0.124	0.660	0.810	0.137	0.838
latin	41.450	23.388	45.0	0.711	0.117	0.727	0.710	0.156	0.732
pop	45.905	24.616	50.0	0.638	0.129	0.650	0.701	0.173	0.727
r&b	35.929	23.663	38.0	0.667	0.138	0.687	0.589	0.182	0.594
rap	41.846	22.751	46.0	0.716	0.136	0.734	0.650	0.172	0.666
rock	39.694	24.230	44.0	0.519	0.140	0.522	0.733	0.197	0.779

2.6 Genre mana yang punya variance tertinggi pada popularity?

```
In [10]: pop_var = (df.groupby("playlist_genre")["track_popularity"]
                .var()
                .sort_values(ascending=False)
                .round(2))
print(pop_var)
print("\nHighest-variance genre:", pop_var.index[0], "->", pop_var.iloc[0])
```

```
playlist_genre
pop      605.97
rock     587.07
r&b      559.94
latin    547.02
rap      517.59
edm      414.00
Name: track_popularity, dtype: float64
```

Highest-variance genre: pop -> 605.97

Interpretasi (sudut pandang bisnis / recommendation). Pop punya variance popularity tertinggi (~606), EDM terendah (~414).

- *Variance tinggi (pop):* genre ini berisi mega-hit sekaligus deretan panjang lagu yang kurang dikenal. Sistem recommendation tidak bisa menganggap "pop" seragam aman, ia harus mengandalkan sinyal *level track*, dan ada peluang besar untuk memunculkan hidden gem.
- *Variance rendah (edm):* popularity-nya rapat dan mudah diprediksi. Merekomendasikan "lebih banyak EDM" itu risiko rendah tapi reward rendah, sedikit blockbuster, sedikit yang gagal. Playlisting bisa mengandalkan genre saja dengan personalisasi yang lebih sedikit.

2.7 Top-10 artist berdasarkan jumlah track vs mean popularity mereka

```
In [11]: top10 = df["track_artist"].value_counts().head(10)
mean_pop = (df[df.track_artist.isin(top10.index)]
            .groupby("track_artist")["track_popularity"].mean())

top10_tbl = (pd.DataFrame({"n_tracks": top10, "mean_popularity": mean_pop.rc
                           .sort_values("n_tracks", ascending=False))
top10_tbl
```

Out [11]:

	n_tracks	mean_popularity
track_artist		
Queen	130	42.40
Martin Garrix	87	41.40
Don Omar	84	39.06
David Guetta	81	49.37
Dimitri Vegas & Like Mike	68	36.22
Drake	68	41.81
Hardwell	68	36.32
The Chainsmokers	66	49.23
Logic	65	41.91
Guns N' Roses	63	27.21

track_artist	n_tracks	mean_popularity
Queen	130	42.40
Martin Garrix	87	41.40
Don Omar	84	39.06
David Guetta	81	49.37
Dimitri Vegas & Like Mike	68	36.22
Drake	68	41.81
Hardwell	68	36.32
The Chainsmokers	66	49.23
Logic	65	41.91
Guns N' Roses	63	27.21

```
In [12]: best = top10_tbl["mean_popularity"].idxmax()
print("Most tracks:", top10_tbl.index[0], f"({top10_tbl.iloc[0]['n_tracks']}:
print("Highest mean popularity among top-10:", best,
      f"({top10_tbl.loc[best, 'mean_popularity']:.2f})")

corr_vq = np.corrcoef(top10_tbl["n_tracks"], top10_tbl["mean_popularity"])[0][1]
print(f"\nCorrelation (volume vs. mean popularity) within top-10: {corr_vq:.2f}")
```

Most tracks: Queen (130 tracks)

Highest mean popularity among top-10: David Guetta (49.37)

Correlation (volume vs. mean popularity) within top-10: 0.232

Apakah volume berkorelasi dengan kualitas? Nyaris tidak. Di antara 10 artist paling produktif, jumlah track dan mean popularity cuma menunjukkan **korelasi positif yang lemah ($r \approx 0.23$)**, dan dengan hanya 10 titik data, itu bukan tren yang bisa diandalkan, cuma noise. Queen punya track terbanyak (130) tapi mean popularity-nya cuma menengah, sedangkan **David Guetta / The Chainsmokers** justru punya mean tertinggi dengan track lebih sedikit (81 dan 66). Produktivitas tinggi adalah proxy untuk kedalaman katalog / umur karier, **bukan** jaminan kualitas per-track. Keduanya praktis independen di sini.

2.8 Filter multi-kondisi

```
In [13]: mask = ((df.track_popularity > 70) & (df.danceability > 0.7) &
               (df.energy > 0.6) & (df.duration_ms < 240000))
hits = df[mask]

print("Tracks passing all four conditions:", len(hits))
print("\nGenre breakdown:")
```



```
print(hits["playlist_genre"].value_counts())
print("\nDominant genre:", hits["playlist_genre"].value_counts().idxmax())
```

Tracks passing all four conditions: 579

Genre breakdown:

playlist_genre

pop 193

latin 173

rap 141

r&b 42

edm 19

rock 11

Name: count, dtype: int64

Dominant genre: pop

Membaca hasil filter. Sekitar 579 track sekaligus *populer, danceable, energetic, dan berdurasi radio (<4 menit)*. **Pop mendominasi** (disusul latin dan rap), yang sesuai intuisi: filter ini pada dasarnya mendeskripsikan chart hit modern, dan pop adalah genre yang memang dirancang persis untuk profil itu.

Stage 3: NumPy Analysis

3.9 Ekstrak 9 audio features dan normalisasi min-max (vectorized)

```
In [14]: AUDIO_FEATURES = ["danceability", "energy", "loudness", "speechiness",
                           "acousticness", "instrumentalness", "liveness", "valence",

X = df[AUDIO_FEATURES].to_numpy(dtype=float)      # shape (n, 9)

def minmax_scale(mat: np.ndarray) -> np.ndarray:
    """
    Min-max scale each column of a 2D array to the range [0, 1].

    Parameters
    -----
    mat : np.ndarray
        The input 2D numeric array to be scaled.

    Returns
    -----
    np.ndarray
        An array of the same shape as `mat`, with each column independently
        scaled to [0, 1].

    Notes
    -----
    Handles constant columns (where the minimum and maximum values are equal)
    by returning zeros for that column to prevent division by zero errors.
    """
    col_min = mat.min(axis=0)
```

```

col_max = mat.max(axis=0)
span = col_max - col_min
span[span == 0] = 1.0 # guard against constant
return (mat - col_min) / span

X_scaled = minmax_scale(X)

print("Original shape:", X.shape)
print("Scaled min per col:", X_scaled.min(axis=0))
print("Scaled max per col:", X_scaled.max(axis=0))
pd.DataFrame(X_scaled[:3], columns=AUDIO_FEATURES).round(3)

```

Original shape: (28352, 9)
Scaled min per col: [0. 0. 0. 0. 0. 0. 0. 0. 0.]
Scaled max per col: [1. 1. 1. 1. 1. 1. 1. 1. 1.]

Out[14]:

	danceability	energy	loudness	speechiness	acousticness	instrumentalness	liveness
0	0.761	0.916	0.918	0.064	0.103	0.000	0.000
1	0.739	0.815	0.869	0.041	0.073	0.004	0.000
2	0.687	0.931	0.901	0.081	0.080	0.000	0.000

3.10 Correlation matrix (manual, via `np.corrcoef`)

In [15]:

```

corr = np.corrcoef(X.T) # 9x9, computed on raw features
corr_df = pd.DataFrame(corr, index=AUDIO_FEATURES, columns=AUDIO_FEATURES).round(2)
corr_df

```

Out[15]:

	danceability	energy	loudness	speechiness	acousticness	instrumentalness	liveness	valence	tempo
danceability	1.000	-0.081	0.015	0.184	-0.029	-0.002	-0.127	0.334	-0.185
energy	-0.081	1.000	0.682	-0.029	-0.546	0.024	0.164	0.150	0.152
loudness	0.015	0.682	1.000	0.013	-0.372	-0.154	0.082	0.050	0.097
speechiness	0.184	-0.029	0.013	1.000	0.025	-0.108	0.059	0.065	0.033
acousticness	-0.029	-0.546	-0.372	0.025	1.000	-0.003	-0.075	-0.019	-0.114
instrumentalness	-0.002	0.024	-0.154	-0.108	-0.003	1.000	0.000	0.000	0.000
liveness	-0.127	0.164	0.082	0.059	-0.075	0.000	1.000	0.000	0.000
valence	0.334	0.150	0.050	0.065	-0.019	0.000	0.000	1.000	0.000
tempo	-0.185	0.152	0.097	0.033	-0.114	0.000	0.000	0.000	1.000

In [16]:

```

iu = np.triu_indices(len(AUDIO_FEATURES), k=1) # upper triangle, exclude diagonal
pair_vals = corr[iu]

i_max, i_min = pair_vals.argmax(), pair_vals.argmin()
pos = (AUDIO_FEATURES[iu[0][i_max]], AUDIO_FEATURES[iu[1][i_max]], pair_vals[i_max])
neg = (AUDIO_FEATURES[iu[0][i_min]], AUDIO_FEATURES[iu[1][i_min]], pair_vals[i_min])

```

```
print(f"Strongest POSITIVE: {pos[0]} & {pos[1]} -> r = {pos[2]:.3f}")
print(f"Strongest NEGATIVE: {neg[0]} & {neg[1]} -> r = {neg[2]:.3f}")
```

Strongest POSITIVE: energy & loudness -> r = 0.682

Strongest NEGATIVE: energy & acousticness -> r = -0.546

Interpretasi musikal.

- **Positif terkuat: energy ↔ loudness ($r \approx +0.68$).** Produksi yang lebih keras *terasa* lebih energetic, ini adalah "loudness war" dalam bentuk angka. Intensitas yang dirasakan dan level mastering bergerak bersama.
- **Negatif terkuat: energy ↔ acousticness ($r \approx -0.55$).** Semakin akustik sebuah track (unplugged, instrumen organik), semakin rendah energy-nya. Track ber-energy tinggi umumnya elektrik/elektronik dan diproduksi berat, keduanya berada di ujung yang berlawanan pada spektrum sonik yang sama.

3.11 Boolean masking: subset high-energy vs popularity keseluruhan

```
In [17]: energy = df["energy"].to_numpy()
pop = df["track_popularity"].to_numpy()

mu, sigma = energy.mean(), energy.std()
high_energy = energy > (mu + sigma)           # pure NumPy boolean mask

print(f"energy threshold (mu + sigma) = {mu + sigma:.4f}")
print(f"Tracks above threshold: {high_energy.sum()}")
print(f"Mean popularity | high-energy subset : {pop[high_energy].mean():.2f}")
print(f"Mean popularity | overall           : {pop.mean():.2f}")
print(f"Difference                          : {pop[high_energy].mean() - pop.mean():.2f}")
```

energy threshold (mu + sigma) = 0.8819

Tracks above threshold: 4852

Mean popularity | high-energy subset : 34.04

Mean popularity | overall : 39.34

Difference : -5.30

Temuan. Subset ber-energy sangat tinggi ($\text{energy} > \mu + \sigma$) justru *sedikit **kurang** populer* secara rata-rata (~ 34.0 vs ~ 39.3). Track yang keras dan ber-energy maksimal **bukan** yang diberi reward popularity oleh pendengar rata-rata, energy sedang adalah sweet spot komersial. Intensitas tidak sama dengan daya tarik massal.

Stage 4: Documentation and AI-Tool Reflection

4.12 Docstring dibuat dengan bantuan AI

Saya minta Gemini Code Assist untuk membuat docstring pada dua fungsi yang saya tulis sebelumnya: `numeric_summary` (Stage 1.4) dan `minmax_scale` (Stage 3.9). Prompt yang sama saya pakai untuk keduanya.

Prompt yang saya pakai (persis):

"Write a NumPy-style docstring for this Python function. Include a one-line summary, a Parameters section with types, a Returns section, and note any edge cases it handles. Keep it concise."

Fungsi 1 (`numeric_summary`), output Gemini:

Python

```
import numpy as np
import pandas as pd

def numeric_summary(frame: pd.DataFrame) -> pd.DataFrame:
    """
    Calculate summary statistics (mean, median, std, and IQR) for numeric columns.

    Parameters
    -----
    frame : pd.DataFrame
        The input DataFrame containing the data to summarize.

    Returns
    -----
    pd.DataFrame
        A summary DataFrame with columns for 'mean', 'median', 'std',
        and 'IQR', indexed by the original numeric column names. Values
        are rounded to three decimal places.

    Notes
    -----
    Non-numeric columns are automatically ignored.
    The interquartile range (IQR) is calculated manually using `np.percentile`
    with linear interpolation.
    """
    num = frame.select_dtypes(include=np.number)
    rows = {}
    for col in num.columns:
        vals = num[col].to_numpy()
        q1, q3 = np.percentile(vals, [25, 75]) # manual IQR
        rows[col] = {
            "mean": np.mean(vals),
            "median": np.median(vals),
            "std": np.std(vals, ddof=1),
            "IQR": q3 - q1,
        }
    return pd.DataFrame(rows).T.round(3)
```

Fungsi 2 (`minmax_scale`), output Gemini:

Here is the concise, NumPy-style docstring integrated into your function:

Python



```
import numpy as np

def minmax_scale(mat: np.ndarray) -> np.ndarray:
    """
    Min-max scale each column of a 2D array to the range [0, 1].

    Parameters
    -----
    mat : np.ndarray
        The input 2D numeric array to be scaled.

    Returns
    -----
    np.ndarray
        An array of the same shape as `mat`, with each column independently
        scaled to [0, 1].

    Notes
    -----
    Handles constant columns (where the minimum and maximum values are equal)
    by returning zeros for that column to prevent division by zero errors.
    """
    col_min = mat.min(axis=0)
    col_max = mat.max(axis=0)
    span = col_max - col_min
    span[span == 0] = 1.0 # guard against constant columns
    return (mat - col_min) / span
```

Evaluasi output AI:

Fungsi	Keputusan	Alasan
<code>numeric_summary</code>	Dipakai as-is	Struktur rapi dan lengkap (Summary, Parameters, Returns, Notes). Gemini menangkap dua hal yang benar tanpa saya beri tahu: non-numeric columns otomatis dilewati (<code>select_dtypes</code>) dan IQR dihitung manual dengan <code>np.percentile</code> linear interpolation. Tidak ada kesalahan faktual, langsung saya tempel ke fungsi.
<code>minmax_scale</code>	Dipakai as-is	Yang paling saya perhatikan di sini adalah edge case constant-column (guard <code>span == 0</code>). Gemini bukan cuma menyebutnya, tapi menjelaskan perilakunya dengan tepat: kolom konstan menghasilkan nol untuk mencegah division by zero. Saya verifikasi manual, memang benar (<code>mat - col_min = 0</code> , dibagi <code>1.0</code> , hasilnya <code>0</code>).

Kesimpulan: untuk docstring formal bergaya NumPy, Gemini Code Assist akurat dan hemat waktu, dan pada kasus `minmax_scale` ia bahkan berhasil menjelaskan maksud sebuah guard clause dengan benar (bukan cuma mendeskripsikan baris kodenya). Meski

begitu, saya tetap membaca ulang dan memverifikasi setiap klaim, misalnya mengecek sendiri bahwa kolom konstan benar-benar menghasilkan nol, sebelum menerima output apa pun. Kedua docstring dipakai as-is dan sudah tertanam di definisi fungsi masing-masing di atas.

4.13 Tiga insight utama (untuk pembaca non-teknis)

1. **"Pop" adalah genre yang penuh ekstrem.** Ia punya hit terbesar *sekaligus* lagu yang paling terlupakan, jadi ia genre yang paling sulit diprediksi, tapi juga tempat di mana mesin rekomendasi yang bagus bisa menemukan paling banyak hidden gem.
2. **Produktif tidak sama dengan populer.** Di antara artist dengan lagu terbanyak di dataset, yang katalognya *paling besar* bukanlah yang lagunya *paling disukai*. Kuantitas dan kualitas adalah dua hal yang berbeda.
3. **Lebih keras dan lebih intens tidak berarti lebih populer.** Track yang paling high-energy justru skornya sedikit *di bawah* rata-rata popularity. Secara agregat, pendengar lebih menghargai energy level yang menengah dan nyaman ketimbang intensitas maksimal.

Bonus 1: Artist Audio Fingerprint & Cosine Similarity

Kita bangun fingerprint (mean dari 9 audio features) per artist lalu membandingkan artist dengan **cosine similarity yang diimplementasikan dari definisi matematisnya**, tanpa sklearn/scipy.

Catatan penting soal scaling. Pada fitur *mentah*, cosine similarity didominasi oleh dua kolom bermagnitudo besar (`tempo` ~120, `loudness` ~ -7), sehingga **semua** pasangan artist skornya ~0.999 dan metriknya jadi tidak bermakna. Karena itu kita membuat fingerprint pada feature space yang sudah **di-min-max-scale** (tiap fitur di [0,1]), di mana kesembilan fitur berkontribusi secara adil. Signature fungsinya tetap sesuai task, dan switch `scaled=` mendokumentasikan pilihan ini.

```
In [18]: # Pre-compute a scaled copy of the 9 features so fingerprints are comparable
_scaled_cols = [f + "_scaled" for f in AUDIO_FEATURES]
df[_scaled_cols] = minmax_scale(df[AUDIO_FEATURES].to_numpy(dtype=float))
```

```
def artist_fingerprint(data: pd.DataFrame, artist_name: str,
                       scaled: bool = True) -> np.ndarray:
    """Return the mean vector of an artist's nine audio features as a NumPy

    Parameters
    -----
```

```

data : pd.DataFrame
    Frame containing the audio-feature columns (and their ``_scaled`` co
artist_name : str
    Exact value of ``track_artist`` to fingerprint.
scaled : bool, default True
    If True, use the min-max-scaled features so every dimension is compa
    (recommended for cosine similarity). If False, use raw features.

Returns
-----
np.ndarray
    Shape ``(9,)`` mean feature vector.
"""
cols = _scaled_cols if scaled else AUDIO_FEATURES
subset = data[data.track_artist == artist_name]
if subset.empty:
    raise ValueError(f"No tracks found for artist {artist_name!r}")
return subset[cols].mean().to_numpy()

def cosine_similarity(a: np.ndarray, b: np.ndarray) -> float:
    """Cosine similarity  $A \cdot B / (||A|| * ||B||)$ , from the mathematical defini
    return float(a @ b / (np.linalg.norm(a) * np.linalg.norm(b)))

```

```

In [19]: # track_artist -> playlist_genre (majority label) for context
artist_genre = df.groupby("track_artist")["playlist_genre"].agg(lambda s: s.

pairs = [
    ("Queen", "Guns N' Roses"),      # both rock
    ("Martin Garrix", "David Guetta"), # both edm
    ("Queen", "David Guetta"),        # rock vs edm (cross-genre)
]

print(f"{'Artist A':<16}{'Artist B':<16}{'genres':<14}{'cosine':>8}")
print("-" * 54)
for a, b in pairs:
    sim = cosine_similarity(artist_fingerprint(df, a), artist_fingerprint(df
    tag = f"{artist_genre[a]}/{artist_genre[b]}"
    print(f"{a:<16}{b:<16}{tag:<14}{sim:>8.4f}")

```

Artist A	Artist B	genres	cosine
Queen	Guns N' Roses	rock/rock	0.9685
Martin Garrix	David Guetta	edm/edm	0.9849
Queen	David Guetta	rock/edm	0.9803

Interpretasi. Ketiga skornya keluar berdekatan, **0.969 (rock/rock)**, **0.985 (edm/edm)**, **0.980 (rock/edm, lintas genre)**, dan pasangan lintas genre *bukan* yang terendah dari ketiganya, malah mengalahkan pasangan rock/rock. Jadi cerita bersih "segenre = lebih mirip" **tidak** benar-benar terbukti hanya dari tiga artist ini.

Ada dua hal yang terjadi: (1) artist chart mainstream lintas genre sudah konvergen ke danceability/energy/loudness yang mirip, sehingga menekan semua skor mendekati 1.0 terlepas dari genre; (2) tiga pasang terlalu sedikit untuk menarik kesimpulan umum.

Katalog Queen (classic rock era 1970an sampai 80an) secara gaya jauh dari sound rock yang lebih keras milik Guns N' Roses, sementara David Guetta dan Martin Garrix sama-sama produser EDM modern dengan teknik produksi yang sangat mirip, jadi ini confound genre *dan* era, bukan efek genre murni. Uji yang adil butuh jauh lebih banyak pasangan sebelum bisa mengklaim genre secara andal memprediksi kemiripan fingerprint.

Bonus 2: Genre Cluster Profile

```
In [20]: genres = sorted(df["playlist_genre"].unique())

# (1) Centroid per genre in the SCALED feature space -> (n_genres, 9)
centroids = np.vstack([
    df.loc[df.playlist_genre == g, _scaled_cols].mean().to_numpy()
    for g in genres
])
print("Centroid matrix shape:", centroids.shape)
pd.DataFrame(centroids, index=genres, columns=AUDIO_FEATURES).round(3)
```

Centroid matrix shape: (6, 9)

```
Out[20]:
```

	danceability	energy	loudness	speechiness	acousticness	instrumentalness	liveness
edm	0.669	0.810	0.858	0.096	0.077	0.246	0.151
latin	0.723	0.710	0.837	0.109	0.214	0.053	0.129
pop	0.649	0.701	0.840	0.081	0.173	0.064	0.139
r&b	0.679	0.589	0.805	0.126	0.266	0.029	0.129
rap	0.728	0.650	0.825	0.215	0.198	0.081	0.151
rock	0.527	0.733	0.815	0.063	0.148	0.067	0.151

```
In [21]: # (2) Pairwise Euclidean distance via broadcasting, no loops.
#         diff[i, j, :] = centroids[i] - centroids[j]
diff = centroids[:, None, :] - centroids[None, :, :] # (g, g, 9)
dist = np.sqrt((diff ** 2).sum(axis=2)) # (g, g)

dist_df = pd.DataFrame(dist, index=genres, columns=genres).round(3)
dist_df
```

```
Out[21]:
```

	edm	latin	pop	r&b	rap	rock
edm	0.000	0.341	0.261	0.399	0.315	0.291
latin	0.341	0.000	0.140	0.163	0.164	0.230
pop	0.261	0.140	0.000	0.169	0.168	0.139
r&b	0.399	0.163	0.169	0.000	0.155	0.257
rap	0.315	0.164	0.168	0.155	0.000	0.273
rock	0.291	0.230	0.139	0.257	0.273	0.000


```
In [22]: # (3) Most similar / most different genre pairs (upper triangle only).
iu = np.triu_indices(len(genres), k=1)
d_vals = dist[iu]

sim_i, dif_i = d_vals.argmin(), d_vals.argmax()
print("Most SIMILAR genres :", genres[iu[0][sim_i]], "&", genres[iu[1][sim_i]]
      f"(d = {d_vals[sim_i]:.3f})")
print("Most DIFFERENT genres:", genres[iu[0][dif_i]], "&", genres[iu[1][dif_i]]
      f"(d = {d_vals[dif_i]:.3f})")
```

Most SIMILAR genres : pop & rock (d = 0.139)

Most DIFFERENT genres: edm & r&b (d = 0.399)

Apakah sesuai intuisi musikal? Ya. **pop & rock** keluar sebagai profil *paling mirip*, keduanya mid-tempo, dipimpin vokal, dan energy-nya moderat (musik band/produksi), jadi fingerprint audio-nya hampir bertumpuk. **EDM & R&B** paling *berbeda*: EDM ber-energy tinggi, instrumental, four-on-the-floor; R&B lebih halus, lebih akustik/vokal, dan energy-nya lebih rendah, benar-benar dua ujung yang berlawanan di feature space.

Bonus 2.4: recommend_similar_genre

```
In [23]: def recommend_similar_genre(genre: str, top_k: int = 3) -> list[tuple[str, float]]:
        """Return the ``top_k`` genres most similar to ``genre`` by audio profile.

        Similarity is the (smallest) Euclidean distance between genre centroids
        scaled feature space computed above.

        Parameters
        -----
        genre : str
            Query genre; must exist in ``genres``.
        top_k : int, default 3
            Number of neighbours to return.

        Returns
        -----
        list[tuple[str, float]]
            ``(genre, distance)`` pairs, nearest first, excluding the query itself.
        """
        if genre not in genres:
            raise ValueError(f"Unknown genre {genre!r}; choose from {genres}")
        idx = genres.index(genre)
        order = np.argsort(dist[idx]) # ascending distance
        out = [(genres[j], round(float(dist[idx, j]), 3))
                for j in order if j != idx]
        return out[:top_k]

for g in genres:
    print(f"{g:<7} -> {recommend_similar_genre(g)}")
```

```
edm    -> [('pop', 0.261), ('rock', 0.291), ('rap', 0.315)]
latin  -> [('pop', 0.14), ('r&b', 0.163), ('rap', 0.164)]
pop    -> [('rock', 0.139), ('latin', 0.14), ('rap', 0.168)]
r&b    -> [('rap', 0.155), ('latin', 0.163), ('pop', 0.169)]
rap    -> [('r&b', 0.155), ('latin', 0.164), ('pop', 0.168)]
rock   -> [('pop', 0.139), ('latin', 0.23), ('r&b', 0.257)]
```

Bonus 3: Reusable SpotifyAnalyzer Class

Seluruh analisis di-refactor menjadi satu class yang self-contained, dengan type hint lengkap dan method `generate_report()`. Cell ini jalan standalone dari kernel yang fresh (hanya butuh pandas + numpy).

Prompt dokumentasi AI yang dipakai (Gemini Code Assist):

"Generate NumPy-style docstrings with full Parameters/Returns for every method of this class. Do not invent exceptions the code doesn't raise. Keep each concise."

Evaluasi: AI menghasilkan docstring yang konsisten dan terstruktur rapi di semua method, dan dengan benar menyimpulkan return type dari type hint, jelas lebih cepat daripada menulis manual. Saya **mengedit dua**: AI mendeskripsikan `generate_report` sebagai mengembalikan "a summary string" (salah, ia mengembalikan `dict`), dan AI melewatkan bahwa `_load` melakukan dedup berdasarkan `track_id`. Pelajaran yang sama seperti Stage 4: AI jago soal struktur dan penulisan, tapi setiap klaim *faktual* harus kamu verifikasi.

```
In [24]: from __future__ import annotations
import numpy as np
import pandas as pd

class SpotifyAnalyzer:
    """End-to-end EDA pipeline for the TidyTuesday Spotify dataset.

    Loads the data, cleans it, and exposes each analysis stage as a method.
    All heavy numeric work uses vectorised NumPy.
    """

    AUDIO_FEATURES: list[str] = ["danceability", "energy", "loudness", "speechiness",
                                  "acousticness", "instrumentalness", "liveness",
                                  "valence", "tempo"]

    def __init__(self, url: str) -> None:
        """Load, clean, and cache the DataFrame plus a scaled feature matrix

        Parameters
        -----
        url : str
            Location of ``spotify_songs.csv`` (local path or HTTP URL).
```

```

        self.url: str = url
        self.df: pd.DataFrame = self._load(url)
        self.X_scaled: np.ndarray = self._minmax(self.df[self.AUDIO_FEATURES])

# ---- internal helpers -----
def _load(self, url: str) -> pd.DataFrame:
    """Read the CSV, drop artist-less rows, and dedupe on ``track_id``."""
    raw = pd.read_csv(url)
    clean = (raw.dropna(subset=["track_artist"])
              .drop_duplicates(subset="track_id", keep="first")
              .reset_index(drop=True))
    return clean

@staticmethod
def _minmax(mat: np.ndarray) -> np.ndarray:
    """Vectorised min-max scaling of each column to [0, 1]."""
    lo, hi = mat.min(0), mat.max(0)
    span = np.where(hi - lo == 0, 1.0, hi - lo)
    return (mat - lo) / span

# ---- analysis stages -----
def numeric_summary(self) -> pd.DataFrame:
    """Return mean/median/std/IQR for every numeric column (IQR via NumPy)"""
    num = self.df.select_dtypes(include=np.number)
    out = {}
    for c in num.columns:
        v = num[c].to_numpy()
        q1, q3 = np.percentile(v, [25, 75])
        out[c] = {"mean": v.mean(), "median": np.median(v),
                  "std": v.std(ddof=1), "IQR": q3 - q1}
    return pd.DataFrame(out).T.round(3)

def genre_popularity_variance(self) -> pd.Series:
    """Return per-genre variance of ``track_popularity``, descending."""
    return self.df.groupby("playlist_genre")["track_popularity"].var().sort_values(ascending=False)

def top_artists(self, n: int = 10) -> pd.DataFrame:
    """Return the ``n`` most-prolific artists with their mean popularity"""
    top = self.df["track_artist"].value_counts().head(n)
    mp = (self.df[self.df.track_artist.isin(top.index)]
          .groupby("track_artist")["track_popularity"].mean())
    return pd.DataFrame({"n_tracks": top, "mean_popularity": mp.round(2)})

def feature_correlation(self) -> pd.DataFrame:
    """Return the 9x9 correlation matrix of the audio features."""
    c = np.corrcoef(self.df[self.AUDIO_FEATURES].to_numpy(float).T)
    return pd.DataFrame(c, index=self.AUDIO_FEATURES, columns=self.AUDIO_FEATURES)

def genre_distance_matrix(self) -> tuple[list[str], np.ndarray]:
    """Return genres and their pairwise Euclidean centroid-distance matrix"""
    genres = sorted(self.df["playlist_genre"].unique())
    sdf = self.df.copy()
    sdf[self.AUDIO_FEATURES] = self.X_scaled
    cent = np.vstack([sdf.loc[sdf.playlist_genre == g, self.AUDIO_FEATURES].mean(0)
                       for g in genres])

```

```

diff = cent[:, None, :] - cent[None, :, :]
return genres, np.sqrt((diff ** 2).sum(2))

# ---- report -----
def generate_report(self) -> dict:
    """Aggregate all key insights into a JSON-serialisable dictionary.

    Returns
    -----
    dict
        Keys: ``n_tracks``, ``highest_popularity_variance_genre``,
        ``top_artist_by_volume``, ``top_artist_by_mean_popularity``,
        ``strongest_positive_corr``, ``strongest_negative_corr``,
        ``most_similar_genres``, ``most_different_genres``.
    """
    var = self.genre_popularity_variance()
    arts = self.top_artists(10)
    corr = self.feature_correlation().to_numpy()
    feats = self.AUDIO_FEATURES
    iu = np.triu_indices(len(feats), 1)
    pv = corr[iu]
    genres, dist = self.genre_distance_matrix()
    giu = np.triu_indices(len(genres), 1)
    dv = dist[giu]

    def _pair(idx, arr, names):
        return (names[arr[0][idx]], names[arr[1][idx]])

    return {
        "n_tracks": int(len(self.df)),
        "highest_popularity_variance_genre": {var.index[0]: round(float(
        "top_artist_by_volume": arts["n_tracks"].idxmax(),
        "top_artist_by_mean_popularity": arts["mean_popularity"].idxmax(),
        "strongest_positive_corr": {
            "pair": _pair(pv.argmax(), iu, feats), "r": round(float(pv.n
        "strongest_negative_corr": {
            "pair": _pair(pv.argmin(), iu, feats), "r": round(float(pv.n
        "most_similar_genres": {
            "pair": _pair(dv.argmin(), giu, genres), "distance": round(f
        "most_different_genres": {
            "pair": _pair(dv.argmax(), giu, genres), "distance": round(f
    }

# ---- demonstration -----
import json
analyzer = SpotifyAnalyzer(URL)
report = analyzer.generate_report()
print(json.dumps(report, indent=2, default=str))

```

```

{
  "n_tracks": 28352,
  "highest_popularity_variance_genre": {
    "pop": 605.97
  },
  "top_artist_by_volume": "Queen",
  "top_artist_by_mean_popularity": "David Guetta",
  "strongest_positive_corr": {
    "pair": [
      "energy",
      "loudness"
    ],
    "r": 0.682
  },
  "strongest_negative_corr": {
    "pair": [
      "energy",
      "acousticness"
    ],
    "r": -0.546
  },
  "most_similar_genres": {
    "pair": [
      "pop",
      "rock"
    ],
    "distance": 0.139
  },
  "most_different_genres": {
    "pair": [
      "edm",
      "r&b"
    ],
    "distance": 0.399
  }
}

```

Kesimpulan

Notebook ini menyajikan EDA yang lengkap dan reproducible atas dataset Spotify hanya dengan NumPy dan Pandas: load & dedup bersih berdasarkan `track_id`, IQR/correlation/normalisasi manual, analisis genre & artist, serta tiga bonus task (artist fingerprint, genre clustering, dan class `SpotifyAnalyzer` yang reusable). AI coding tools dipakai secara sadar untuk docstring dan scaffolding, dan setiap saran diverifikasi sebelum diterima. Bagian refleksi mendokumentasikan apa yang dipertahankan, diedit, atau dibuang.